

ChromaDB — Full Documentation (Beginner to Advanced)

1. Overview

ChromaDB is an open-source **vector database** designed for: - Semantic search - RAG (Retrieval Augmented Generation) - AI assistants with memory - Embedding storage & similarity search

ChromaDB stores **vectors**, **documents**, and **metadata**, enabling fast information retrieval based on meaning.

2. Why ChromaDB? (Problem → Solution)

✗ Traditional Search Problem

- Only works with **keywords**
- Misses meaning
- Cannot use embeddings

✓ ChromaDB Solution

- Stores embeddings (vectors)
 - Performs **semantic search**
 - Integrates with LLMs (OpenAI, HuggingFace, LangChain)
 - Extremely fast, local, persistent
-

3. Installation

```
pip install chromadb
```

For persistence (local DB storage):

```
chromadb.PersistentClient(path=".chroma_db")
```

4. Project Folder Structure

```
my_rag_app/  
|  
|--- data/
```

```
|- chroma/           # ChromaDB persistent storage  
|  
+- src/  
|   +- ingest.py    # Add documents + embed  
|   +- query.py     # Query the DB  
|   +- utils.py      # Helpers  
  
- requirements.txt
```

5. Core API Flow (Explained)

```
Raw Text → Embedding → Stored in ChromaDB → Semantic Search → Relevant Output
```

5.1 Steps

1. Initialize ChromaDB client
2. Create / Load a collection
3. Insert documents
4. Query with semantic search
5. (Optional) Update/Delete data

6. Basic Usage Examples

6.1 Create / Load Database

```
import chromadb  
client = chromadb.PersistentClient(path="data/chroma")  
collection = client.get_or_create_collection("my_docs")
```

6.2 Add Documents

```
docs = [  
    "Python is a programming language.",  
    "Django is a Python web framework."  
]  
  
collection.add(  
    ids=["1", "2"],  
    documents=docs  
)
```

6.3 Query Documents

```
result = collection.query(  
    query_texts=["what is django"],  
    n_results=2  
)  
print(result)
```

Returns: - Similar documents - Their distances - Their metadata (if present)

7. Metadata Support

Store extra information with each document.

```
collection.add(  
    ids=["3"],  
    documents=["Flask is a micro web framework."],  
    metadatas=[{"category": "python"}]  
)
```

Query with metadata filter:

```
collection.query(  
    query_texts=["framework"],  
    where={"category": "python"}  
)
```

8. Updating and Deleting Data

Update

```
collection.update(  
    ids=["1"],  
    documents=["Python is a widely used programming language."]  
)
```

Delete

```
collection.delete(ids=["2"])
```

9. Embedding Models Integration

You can plug in: - OpenAI embedding API - HuggingFace SentenceTransformers - LangChain embedding wrappers

Example (HuggingFace)

```
from sentence_transformers import SentenceTransformer
model = SentenceTransformer('all-MiniLM-L6-v2')

vector = model.encode(["Hello world"])
```

10. ChromaDB + LangChain Flow

```
Documents → Embeddings → Chroma Vector Store → LangChain Retriever → LLM
```

Example

```
from langchain.vectorstores import Chroma
from langchain.embeddings import OpenAIEMBEDDINGS

emb = OpenAIEMBEDDINGS()
vector_db = Chroma(
    collection_name="my_docs",
    embedding_function=emb,
    persist_directory="data/chroma"
)
```

11. RAG Example (Full Flow)

```
query = "Explain Django in short"
retriever = vector_db.as_retriever()
context = retriever.get_relevant_documents(query)

print(context)
```

12. Best Practices

- Use **persistent mode** for production
- Use **metadata** for filtering

- Chunk long documents before embedding
 - Use batch insertion for speed
-


13. Limitations


Limitation	Notes
Not a full DB	Only vector storage
Not distributed	But can scale locally
Requires embeddings separately	You choose the embedding model

14. When to Use ChromaDB

- AI chatbots with memory
 - Semantic search engines
 - RAG systems
 - Document Q&A
 - Code search
-

15. Architecture Diagram (Text-Based)





16. One-Page Summary

ChromaDB = Fast vector database for storing embeddings + semantic search. Perfect for RAG and LLM apps.

If you want, I can also add: - Architecture diagram - Advanced performance tuning - A complete RAG mini-project using ChromaDB - PDF export of this documentation